

COMP3220 — Document Processing and the Semantic Web

Week 05 Lecture 1: Processing Text Sequences

Diego Mollá

Department of Computer Science
Macquarie University

COMP3220 2020H1

Programme

1 Word Embeddings

- Challenges of Text for Machine Learning
- Word Embeddings

2 Text Sequences

- Modelling Text Sequences
- Sequence Labelling

Reading

- Deep Learning book, chapter 6.
- Understanding LSTM Networks, <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>.

Programme

1 Word Embeddings

- Challenges of Text for Machine Learning
- Word Embeddings

2 Text Sequences

- Modelling Text Sequences
- Sequence Labelling

Programme

1 Word Embeddings

- Challenges of Text for Machine Learning
- Word Embeddings

2 Text Sequences

- Modelling Text Sequences
- Sequence Labelling

Words as Arbitrary Symbols

- Words are encoded as arbitrary symbols.
- Within one language there is no clear correspondence between a word symbol and its meaning.
 - “dig” vs. “dog”
 - “car” vs. “automobile”
- Different languages may use different representations of the same word.



https://en.wikipedia.org/wiki/File:Hello_in_different_languages_word_cloud.jpeg

Ambiguities Everywhere

Language features ambiguity at multiple levels.


Lexical Ambiguity

Example from Google's dictionary:

- bank (n): the land alongside or sloping down a river or lake.
- bank (n): financial establishment that uses money deposited by customers for investment, . . .
- bank (v): form in to a mass or mound.
- bank (v): build (a road, railway, or sports track) higher at the outer edge of a bend to facilitate fast cornering.
- . . .

So many words!

- Any language features a large number of distinct words.
- New words are coined.
- Words change their use in time.
- There are also names, numbers, dates... an infinite number.



<https://trends.google.com>

Long-distance Dependencies

- Sentences are sequences of words.
- Words close in the sentence are often related.
- But sometimes there are relations between words far apart.

grammatical: “The man living upstairs is very cheerful”

“The people living upstairs are very cheerful”

reference: “I bought a book from the bookshp and I liked it”

knowledge: “I was born in France and I speak fluent French”

Programme

1 Word Embeddings


- Challenges of Text for Machine Learning
- Word Embeddings

2 Text Sequences


- Modelling Text Sequences
- Sequence Labelling

Word Embeddings


- First introduced in 2013, nowadays is one of the most common ingredients in text processing systems.
- Word embeddings squarely aim at addressing the issue of representing words as continuous vectors of integers.
- Words with similar context are mapped to similar vectors.
- Embeddings are learnt using large, unlabelled training data.



Male-Female



Verb tense




Country-Capital

<https://www.tensorflow.org/tutorials/representation/word2vec>

One-hot vs. word embeddings


One-hot

- Sparse
- Binary values (typically)
- High-dimensional
- Hard-coded



Word embeddings

- Dense
- Continuous values
- Lower-dimensional
- Learned from data



Two Ways to Obtain Word Embeddings

- 1 Learn the word embeddings jointly with the task you care about (e.g. document classification).
- 2 Use pre-trained word embeddings.


Learning Word Embeddings

- You can add a dense layer as the first layer of your network and let the system learn the optimal weights.
- This approach is so useful and common that many deep learning frameworks define an “embedding” layer that facilitates this.
- The input to the “embedding” layer is the word index.
- The output is the word embedding.

Embedding Layer as a Dense Layer

The input of the dense layer is the one-hot encoding of the word


A Dense Layer



Embedding Layer in Keras

The input of a Keras embedding layer is a word index which will be internally converted into its one-hot representation.

An Embedding Layer



Using pre-trained word embeddings

The Problem: Data Sparsity

- Sometimes we have so little training data that many words are poorly represented.
- Often, words in the training data do not occur in the test data.
- For these words we would not be able to learn the embeddings.

A Solution: Pre-training

- Several people have computed word embeddings for large vocabularies using large data sets.
- We can then use these pre-trained embeddings to map from the word index to the word embedding.

Using Word Embeddings in Keras

- The following notebook is based on the jupyter notebooks provided by the Deep Learning book: <https://github.com/fchollet/deep-learning-with-python-notebooks>
 - Using word embeddings.
- The notebook illustrates how you can use an embeddings layer for text classification, and how to load pre-trained word embeddings.
- This notebook is important because it also illustrates Keras' text tokenisation techniques.

Programme

1 Word Embeddings

- Challenges of Text for Machine Learning
- Word Embeddings

2 Text Sequences

- Modelling Text Sequences
- Sequence Labelling

Programme

1 Word Embeddings

- Challenges of Text for Machine Learning
- Word Embeddings

2 Text Sequences

- Modelling Text Sequences
- Sequence Labelling

Handling Text Sequences

- A document is a sequence of words.
- Many document representations are based on a bag-of-words approach.
 - Word order is ignored.
 - The context around a word is ignored.
- Even word embeddings ignore word order.


Why context matters

“I can₁ kick the can₂”

- The meaning of “can₁” is different from that of “can₂”.
- “can₁” and “can₂” should have different word embeddings.
- We can tell the meaning because of the context:
 - “I can kick ...”
 - “...kick the can”


Recurrent Neural Networks

- A **Recurrent Neural Network** (RNN) is designed to process sequences.
- A RNN is a neural network that is composed of RNN cells.
- Each RNN cell takes as input two pieces of information:
 - 1 A vector representing an item x_i in the sequence.
 - 2 The state resulting from processing the previous items.
- The output of the RNN cell is a state that can be fed to the next cell in the sequence.
- All cells in an RNN chain share the same parameters.
- In a sense, we can say that an RNN cell is the same for all words in the sequence, but now context also matters.



A Simple Recurrent Neural Networks

- A simple RNN cell (“vanilla RNN”) has just a dense layer with an activation function (hyperbolic tangent, or “tanh” in the drawing below).
- Vanilla RNN cells have been used since 1990s.




<https://colah.github.io/posts/2015-08-Understanding-LSTMs/>

LSTMs and GRUs

- Vanilla RNN cells are still too simple and they do not handle long-distance dependencies easily.
- More complex RNN cells have been designed specifically to address this issue.
- Current most popular RNN cells are:

LSTM Long Short Term Memory (picture).

GRU Gated Recurrent Unit; a more recent, simpler cell.



<https://colah.github.io/posts/2015-08-Understanding-LSTMs/>

RNNs in Practice

- Most deep learning frameworks include special layers for RNNs.
- When you use an RNN layer, you have the option to specify the type of RNN cell.
- You often have the option to use the state of the last cell, or the state of all cells.

Recurrent Neural Networks in Keras

The following notebook is based on the jupyter notebooks provided by the Deep Learning book: <https://github.com/fchollet/deep-learning-with-python-notebooks>

- Understanding Recurrent Neural Networks.

The notebook illustrates how you can use an embeddings layer for text classification, and how to load pre-trained word embeddings.

Programme

1 Word Embeddings

- Challenges of Text for Machine Learning
- Word Embeddings

2 Text Sequences

- Modelling Text Sequences
- Sequence Labelling

What is Sequence Labelling?

- A **sequence labelling** problem is one where:
 - the input consists of a sequence $\mathbf{X} = (X_1, \dots, X_n)$, and
 - the output consists of a sequence $\mathbf{Y} = (Y_1, \dots, Y_n)$ of labels, where:
 - Y_i is the label for element X_i
- Example: Part-of-speech tagging

$$\begin{pmatrix} \mathbf{Y} \\ \mathbf{X} \end{pmatrix} = \begin{pmatrix} \text{Verb,} & \text{Determiner,} & \text{Noun} \\ \text{spread,} & \text{the,} & \text{butter} \end{pmatrix}$$

- Example: Spelling correction

$$\begin{pmatrix} \mathbf{Y} \\ \mathbf{X} \end{pmatrix} = \begin{pmatrix} \text{write,} & \text{a,} & \text{book} \\ \text{rite,} & \text{a,} & \text{buk} \end{pmatrix}$$

Other applications of sequence labelling

- **Named entity recognition** and classification (NER) involves finding the named entities in a text and identifying what type of entity they are (e.g., person, location, corporation, dates, etc.).
- **Speech transcription** can be seen as a sequence labelling task:
 - The input $\mathbf{X} = (X_1, \dots, X_n)$ is a sequence of **acoustic frames** X_i , where X_i is a set of features extracted from a 50msec window of the speech signal.
 - The output \mathbf{Y} is a sequence of words (the transcript of the speech signal).
- **Financial applications** of sequence labelling:
 - Identifying trends in price movements.
- **Biological applications** of sequence labelling:
 - Gene-finding in DNA or RNA sequences.

Sequence Labelling as Classification I

Can we just use a standard classifier?

- Standard classifiers (such as those we have seen so far) assume **independence between samples**:
 - The probability of the label assigned to sample i is independent to the probability of the label assigned to sample j .
- But in sequence labelling there is interdependence between the labels of different samples.

Modelling Context

Classifier with context features

- A (crude) approach to model interdependence between samples is to add context features.
- For example, we can use features based on previous words and following words.
- We can even incorporate the label of the previous word as a feature.
- But it is not so easy to incorporate the label of **both** the previous word and the following word.

Using Recurrent Neural Networks for Sequence Labelling I

- We have seen how RNNs can be used to classify documents.
- Similarly, we can use RNNs to classify sequences of words.
- Keras provides 'TimeDistributed', which assigns a copy of a layer to the output of each recurrent cell.

```
model = Sequential()  
model.add(Embedding(max_words, 32))  
model.add(LSTM(32, return_sequences=True))  
model.add(TimeDistributed(Dense(num_tags,  
                                activation='softmax')))  
  
model.compile(optimizer='Adam',  
              loss='categorical_crossentropy',  
              metrics=['acc'])
```

Using Recurrent Neural Networks for Sequence Labelling II

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, None, 32)	320000
lstm_1 (LSTM)	(None, None, 32)	8320
time_distributed_1 (TimeDist	(None, None, 20)	660

Sequence Labelling in Keras

- The following notebook is based on nlpforhackers' post:
<https://nlpforhackers.io/lstm-pos-tagger-keras/>
 - Building a Part of Speech Tagger with Keras
- We will look at how to build a Part of Speech tagger using a LSTM layer.
- The notebook uses NLTK's treebank corpus, which contains annotations with the parts of speech of the words.

Final Note: Contextualised Word Embeddings!



Recent research devised a way to combine RNN and word embeddings to produce context-dependent word embeddings. The resulting systems are beating state of the art in many applications!



<http://jalammar.github.io/illustrated-bert/>

Take-home Messages

- 1 Explain some of the fundamental challenges that plain text represents to machine learning.
- 2 Apply word embeddings in deep learning.
- 3 Use recurrent neural networks for text classification.
- 4 Comment on the issues of sequence labelling.

What's Next

Week 6

- Generating text.
- Reading: Deep Learning book, chapter 8.1